

Raven: Scheduling Virtual Machine Migration during Data Center Upgrades with Reinforcement Learning

Chen Ying
Baochun Li
Xiaodi Ke
Lei Guo

TODAY'S AGENDA

- 1 Background & Context
- 2 Brief Deep RL Summary
- 3 Raven:
The Scheduler
- 4 Results
- 5 Conclusions
- 6 How does this relate?



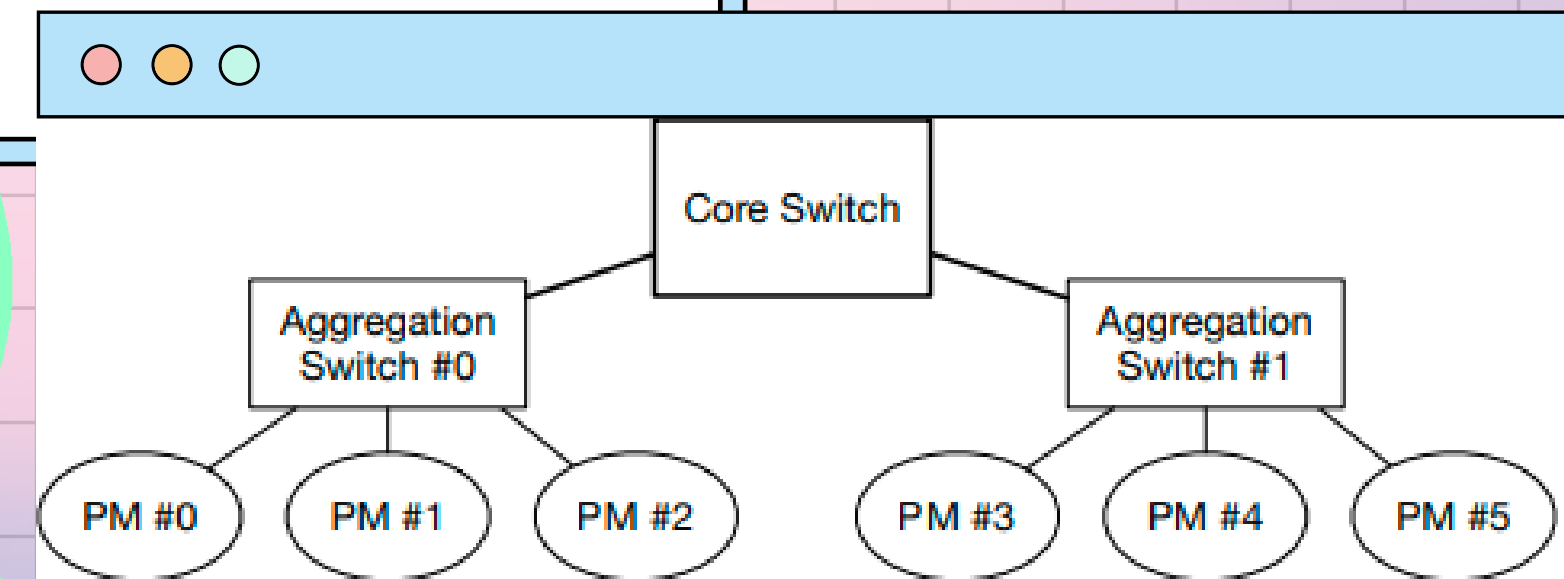
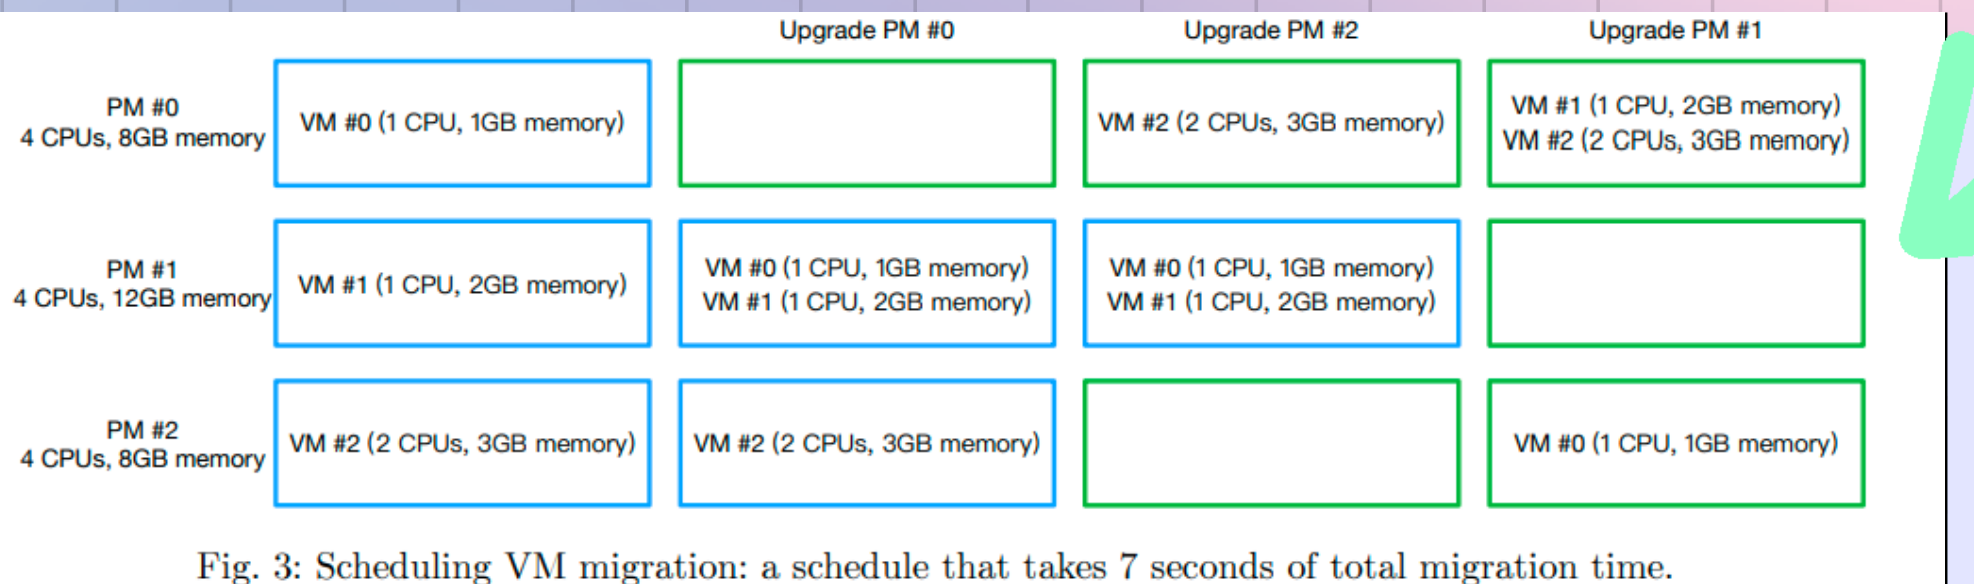
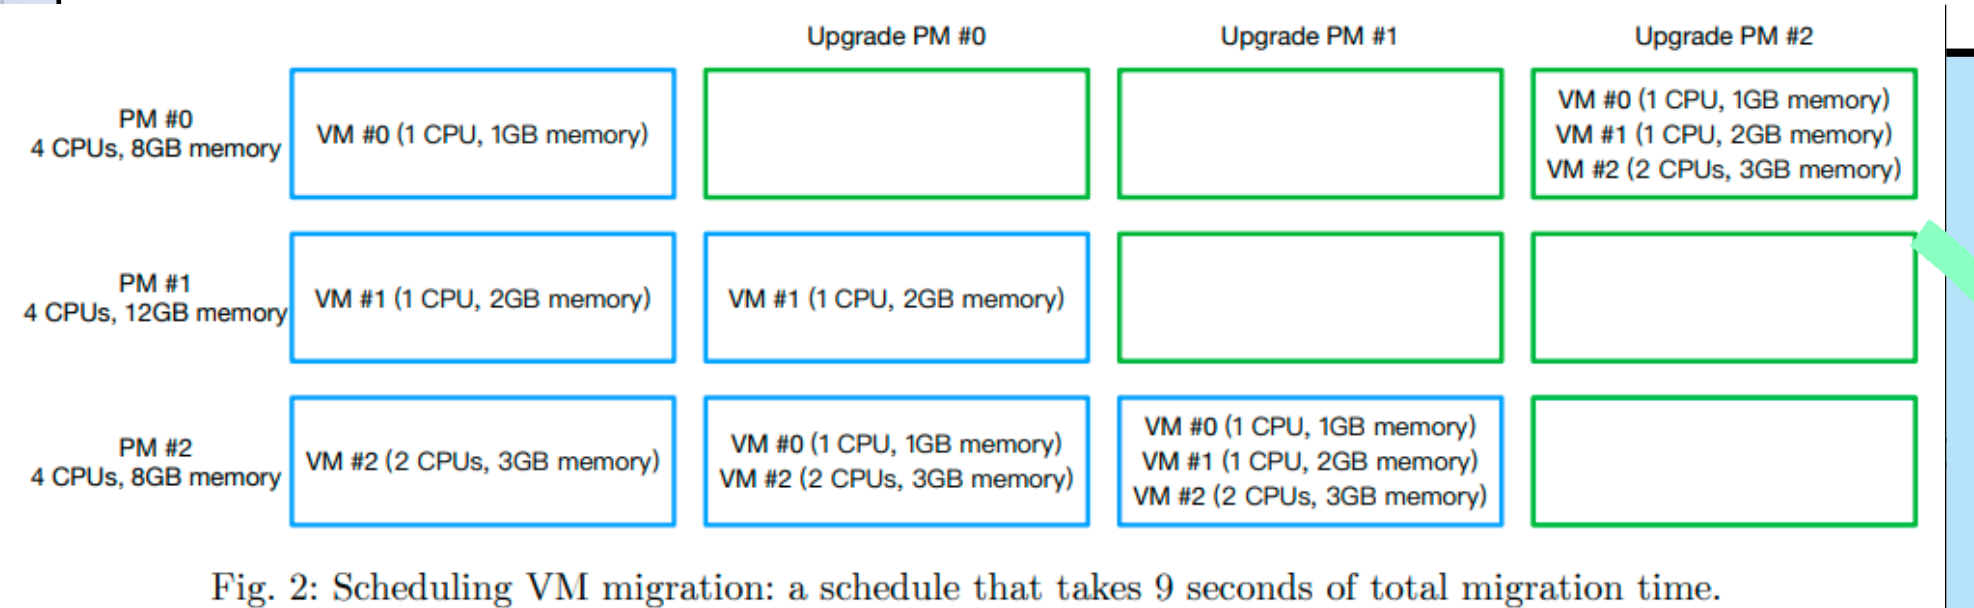
BACKGROUND & CONTEXT

- Common for modern data centers to require maintenance upgrades for physical machines(PMs)
 - migrate virtual machines (VMs)
 - reduce downtime and/or disruptions
 - migrating images takes the longest
- Must carefully select destination PM and schedule the VM migration
- Not much related works or within 'normal' situations
 - network topology and link capacities would be initially unknown

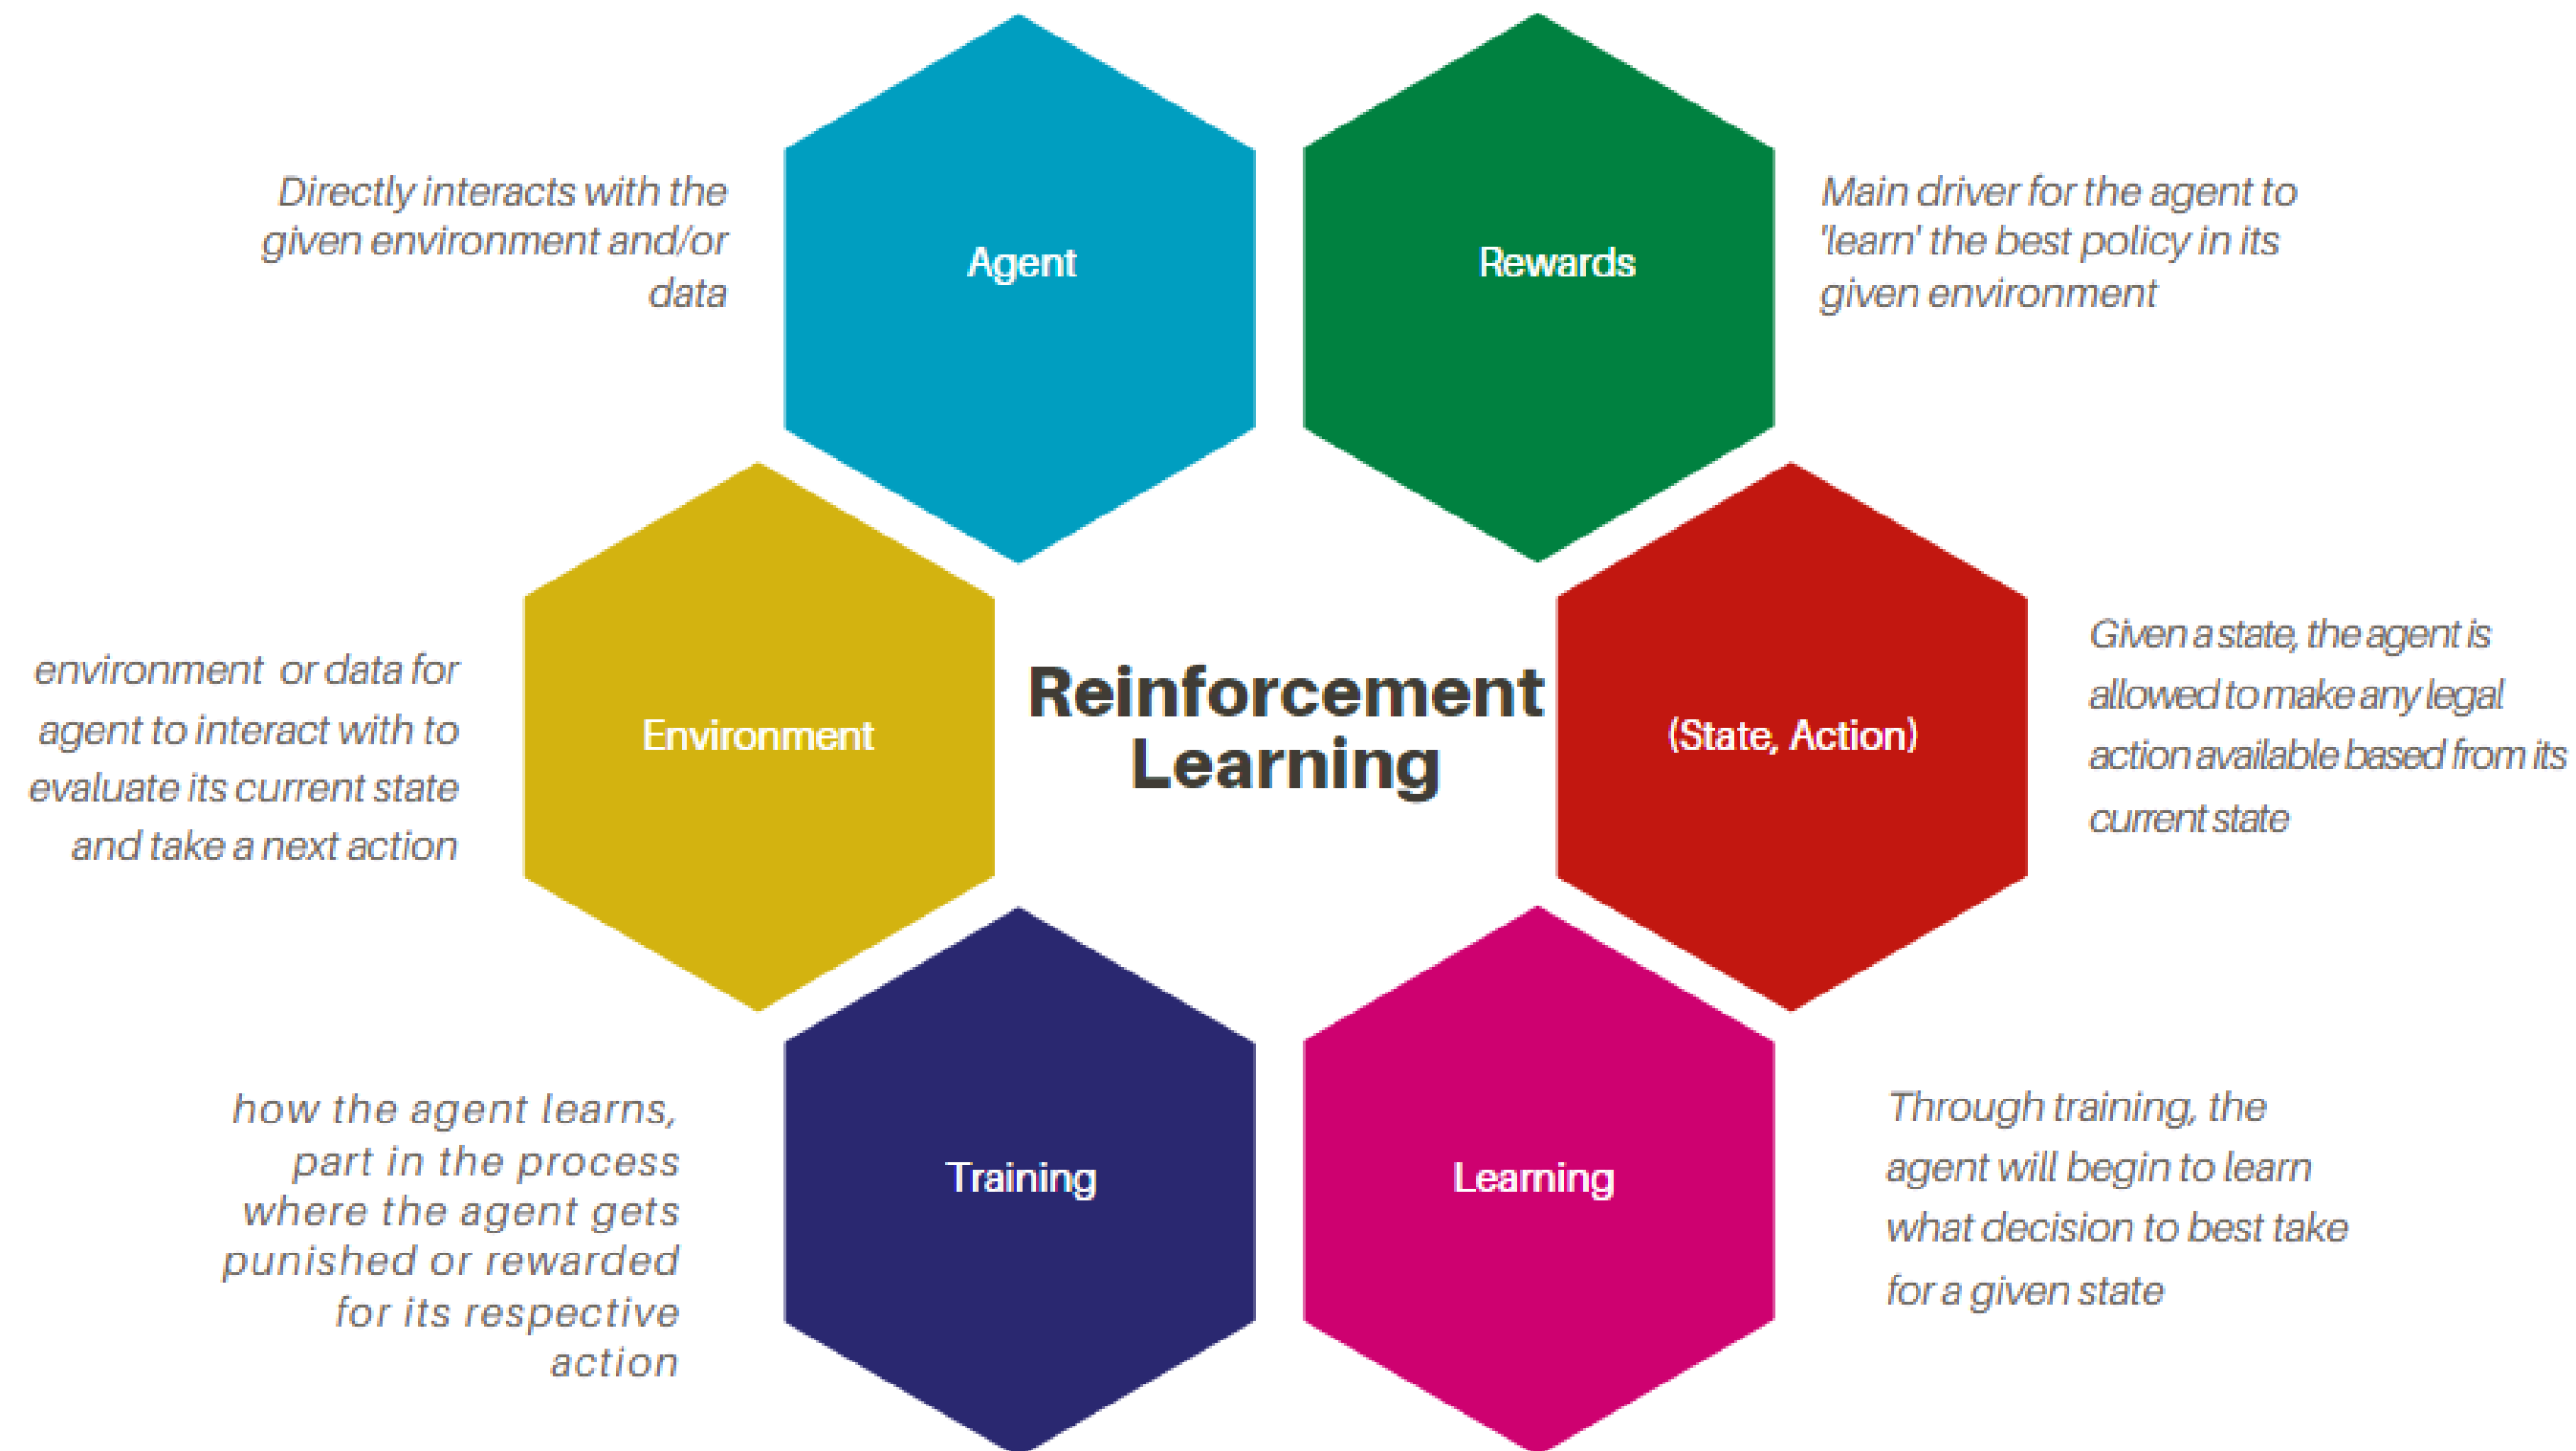


Problem to Solve

Can we leverage deep reinforcement learning to schedule VM migration to ensure the lowest total migration time to upgrade the PMs? And can this be done without prior knowledge of the network topology and network link capacities



BRIEF REINFORCEMENT LEARNING SUMMARY



Extra features

- Fully connected neural network
 - adjustable policy $\pi(a|s; \theta)$ and parameters
- Cross-Entropy method is used in calculations to find optimal policy $\pi(a|s; \theta^*)$

Parameter for -> sampling (3)

$$\hat{v} = \arg \max_v \frac{1}{N} \sum_{n \in [N]} \mathbf{1}_{\{R(x_n) \geq \xi\}} \frac{f(x_n; u)}{f(x_n; w)} \log f(x_n; v), \quad (3)$$

$$\hat{\theta}_k = \arg \max_{\theta_k} \sum_{n \in [N]} \mathbf{1}_{\{R(x_n) \geq \xi_k\}} \left(\sum_{a_t, s_t \in x_n} \pi(a_t | s_t; \theta_k) \right), \quad (4)$$

<- Parameter estimator at iteration 'k' (4)

It's reinforcement learning with a few extra features..

RAVEN?

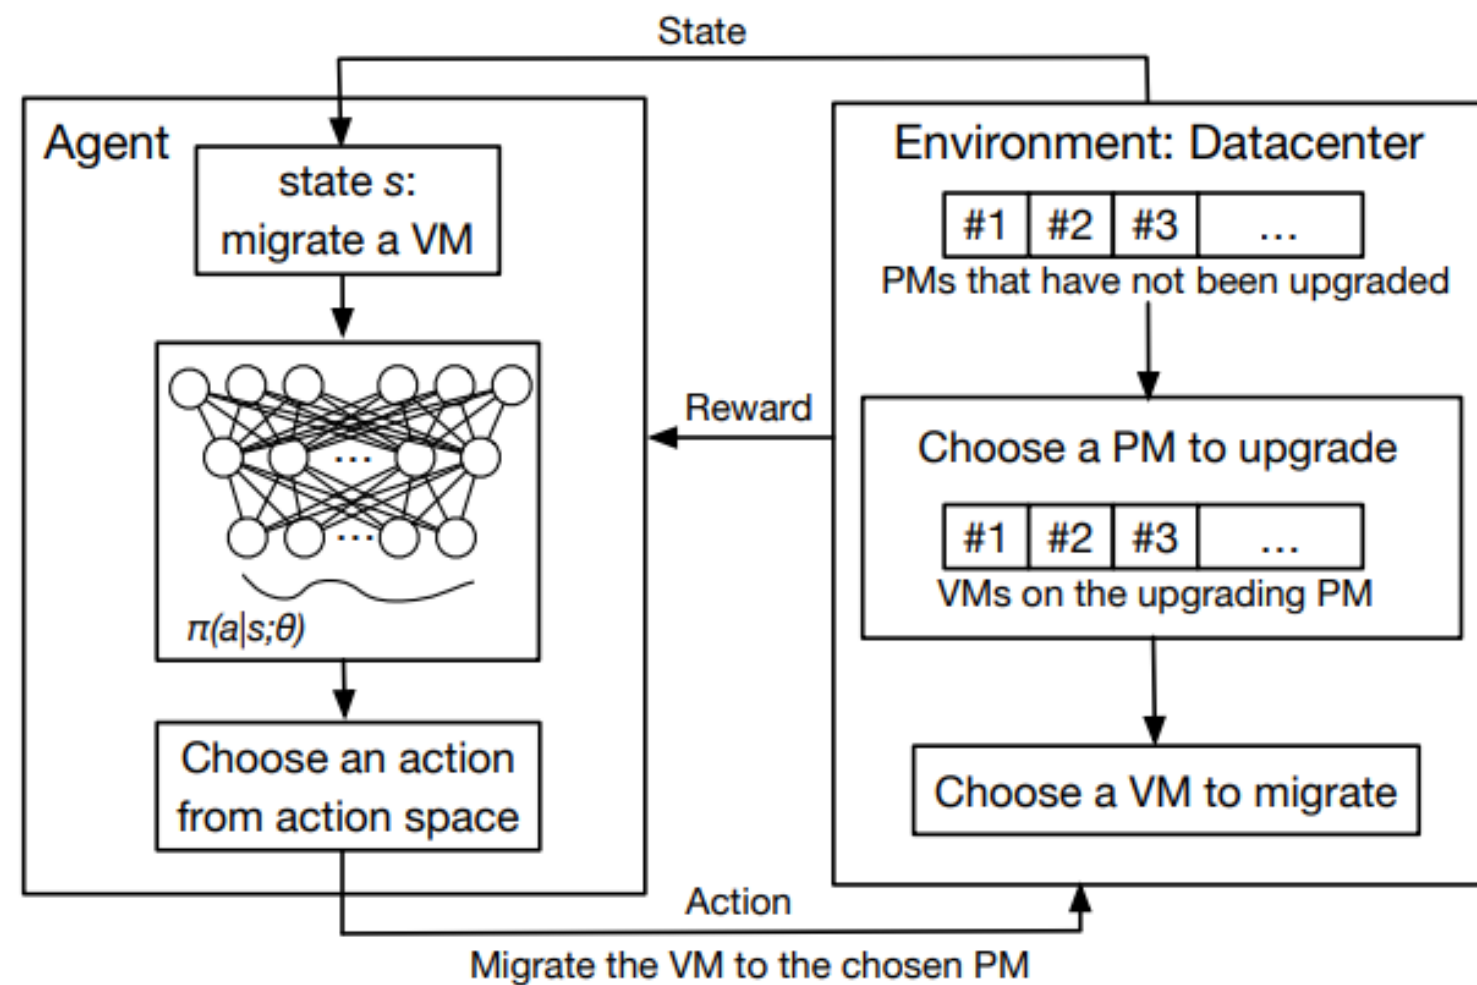


Fig. 4: The architecture of *Raven*.

State of the environment:

$$s_t = \{s_{t1}, s_{t2}, \dots, s_{tJ}, v_t^{\text{cpu}}, v_t^{\text{mem}}, v_t^{\text{pm id}}\},$$

Episode: Finish upgrading all PMs
Time step

Start: Pick PM that needs to be upgraded

Per timestep: VM is migrated

Action: destination PM index

State Space: $s_{tj} = \{s_{tj}^{\text{status}}, s_{tj}^{\text{total cpu}}, s_{tj}^{\text{total mem}}, s_{tj}^{\text{used cpu}}, s_{tj}^{\text{used mem}}\},$

Reward: lower total migration time

RESULTS



Experiment was done on varying network topologies that were either 2-layer and 3-layer settings with 16 episodes for each setting




Fig. 5: The learning curve of *Raven* in datacenter with 9 physical machines and 30 virtual machines.




Fig. 6: The learning curve of *Raven* in datacenter with 10 physical machines and 40 virtual machines.

RESULTS CTD




Fig. 8: The learning curve of *Raven* in datacenter with 260 physical machines and 520 virtual machines.




Fig. 9: The learning curve of *Raven* in datacenter with 260 physical machines and 1158 virtual machines.

RESULTS CTD



Table 1: Total migration time within different datacenters.

Datacenter Setting			Total Migration Time		
Number of PMs	Number of VMs	Number of Aggregation Switches	Min-DIFF	Heuristic	<i>Raven</i>
50	100	0	35.56 (11%)	57.50 (45%)	31.50
50	100	2	130.5 (27%)	110.0 (14%)	94.00
50	150	0	55.50 (31%)	103.00 (63%)	38.00
50	150	2	168.00 (16%)	159.00 (11%)	140.50
100	200	0	43.50 (28%)	91.50 (66%)	31.00
100	200	2	227.00 (35%)	221.50 (34%)	146.00
100	300	0	65.50 (2%)	192.50 (66%)	63.99
100	300	2	339.00 (15%)	378.00 (24%)	286.00
260	520	6	791.00 (58%)	510.00 (34%)	332.14
260	520	7	681.00 (35%)	445.78 (1%)	440.15
260	520	8	651.25 (58%)	440.00 (39%)	268.00
260	1158	6	944.31 (7%)	924.42 (5%)	869.97
260	1158	7	903.00 (10%)	910.50 (11%)	805.85
260	1158	8	918.50 (15%)	870.95 (11%)	774.23

RESULTS CTD



Same datacenter settings but with 10 different VM mappings

Table 2: Average total migration time within different datacenter.

Datacenter Setting			Average Total Migration Time		
Number of PMs	Number of VMs	Number of Aggregation Switches	Min-DIFF	Heuristic	<i>Raven</i>
50	100	0	36.20 (0%)	57.55 (36%)	36.39
50	100	2	117.00 (1%)	107.38 (-7%)	115.13
50	150	0	47.25 (-1%)	94.50 (49%)	47.94
50	150	2	172.95 (9%)	165.83 (5%)	157.00
100	200	0	53.05 (17%)	100.16 (56%)	43.54
100	200	2	248.24 (15%)	214.80 (2%)	209.14
100	300	0	66.80 (7%)	156.65 (60%)	61.96
100	300	2	335.95 (12%)	299.65 (1%)	296.32
260	520	6	725.55 (46%)	451.52 (14%)	388.94
260	520	7	667.70 (35%)	465.76 (7%)	433.17
260	520	8	622.65 (45%)	420.08 (19%)	338.65
260	1158	6	971.67 (12%)	1020.25 (17%)	846.06
260	1158	7	912.01 (14%)	987.62 (20%)	782.55
260	1158	8	836.54 (10%)	969.42 (22%)	748.93

CONCLUSIONS



Needed Improvements:

- consider live VM migration
- struggles to converge under certain settings

Overall


- ✓ closer to a real-life use case scenario
- ✓ does better in larger network topologies
- ✓ still has the shortest times overall compared to:
 - Min-DIFF
 - Heuristic evaluation

HOW DOES THIS RELATE?



- ✓ Deep reinforcement learning looks more promising to pursue
 - large fat tree, harder to maintain a traditional Q Table

- ✓ Similar unique example, with no obvious starting event
 - good example to 'translate'



Works Cited

**Ying, C., Li, B., Xiaodi, K.,
& Guo, L. (2020). Raven :
Scheduling Virtual Machine
Migration During Datacenter
Upgrades with Reinforcement
Learning. *Mobile Networks and
Applications*, 1-12.**

○ ○ ○ ○

*Thank
you!*

**Any
questions,
comments,
and/or
concerns?**